

SECTION 03 · DATABASE

PostgreSQL Administration & SQL

Roles, SQL fluency, indexes, EXPLAIN, backups for Odoo databases

AGENDA

What this section teaches

Setup

Install, roles, grants, config files

SQL

DDL, CRUD, joins, aggregations

Performance

Indexes, EXPLAIN, locks

Ops

Dump/restore and Odoo DB habits

16 lessons · explanations · examples · labs

METHOD

How to teach and learn this deck

- Each lesson starts with objectives, then teaching narrative
- Code/command slides are meant to be typed live or in labs
- Pitfall slides capture production failure modes
- Use the matching documentation part for full prose depth

Facilitation: These are lecture materials—pause for learner typing time after every example.

LESSONS DB01–DB05

PostgreSQL Administration & SQL

Teaching block 1

DB01 · OBJECTIVES

Install PostgreSQL on Ubuntu

- Install PostgreSQL packages on Ubuntu in a repeatable way.
- Understand the system service, default cluster, and postgres OS user.
- Verify readiness before connecting Odoo to the server.

Lab target: Install PostgreSQL on a fresh Ubuntu machine, confirm the service is active, and capture the output of `SELECT version()`.

DB01 · TEACHING 1

Install PostgreSQL on Ubuntu

A clean PostgreSQL installation starts with knowing which package source you trust. Ubuntu ships PostgreSQL packages that are stable and easy to patch, while the upstream PostgreSQL APT repository gives access to newer major versions sooner.

On Ubuntu, PostgreSQL is managed as a systemd service and normally creates a default cluster during installation. The cluster is a data directory plus configuration files, not just a running process.

DB01 · TEACHING 2

Install PostgreSQL on Ubuntu (continued)

The postgres Linux user is intentionally separate from your application user. It owns the database files and can connect locally as the postgres database superuser through peer authentication.

For Odoo training and development, keep the database server local unless there is a clear operational reason to separate it. Localhost reduces network complexity and makes `pg_hba.conf` easier to reason about.

DB01 · TEACHING 3

Install PostgreSQL on Ubuntu (continued)

After installation, verify the service state, listening port, and psql connection before installing Odoo. Debugging Odoo connection errors is much faster when the database layer is already proven healthy.

Production installations should also consider disk layout, backups, monitoring, and upgrade policy before real data is loaded.

PostgreSQL is easy to install, but moving a busy Odoo database later is not always simple.

DB01 · KEY POINTS

What to remember

- Use apt for predictable installs and security updates.
- The default service is usually postgresql, with one or more versioned clusters behind it.
- The postgres OS account is for database administration, not for running Odoo.
- Always verify the server with psql before blaming Odoo.

DB01 · EXAMPLE 1

Ubuntu package installation

Ubuntu package installation

This installs PostgreSQL, starts the service, and proves that the local postgres administrator can execute SQL.

```
sudo apt update
sudo apt install -y postgresql postgresql-contrib
sudo systemctl enable --now postgresql
systemctl status postgresql --no-pager
sudo -u postgres psql -c "SELECT version();"
```

DB01 · PITFALLS

Common mistakes

- Installing PostgreSQL but never checking whether the service is actually running.
- Running Odoo as the postgres Linux user instead of creating a dedicated Odoo database role.
- Mixing Ubuntu packages and upstream packages without documenting the version source.
- Opening PostgreSQL to the network before authentication rules and firewall rules are understood.

Remember: Validate fixes in a disposable database before production.

DB01 · EXTRA

Install checks

Check	Command	Healthy result
Service	systemctl status postgresql	active (running)
Socket	ss -ltnp rg 5432	Local listener exists
SQL	sudo -u postgres psql -c 'SELECT 1;'	Returns one row

DB02 · OBJECTIVES

psql CLI, Users, Roles, and Grants

- Connect with psql using local and password authentication.
- Create roles suitable for Odoo development.
- Grant only the privileges an application role needs.

Lab target: Create a role named `odoo17_lab` with `CREATEDB`, connect with psql over `127.0.0.1`, and list all roles with `\du`.

DB02 · TEACHING 1

psql CLI, Users, Roles, and Grants

psql is the fastest diagnostic tool for PostgreSQL because it removes the application layer from the problem. A developer who can inspect roles, databases, schemas, and locks from psql can solve most Odoo database issues directly.

PostgreSQL uses roles for both users and groups. A role can log in when it has LOGIN, and it can own objects, inherit privileges, or be granted to other roles depending on how you design access.

DB02 · TEACHING 2

psql CLI, Users, Roles, and Grants (continued)

Odoo usually connects with one role that owns or can create its databases in development. In production, many teams restrict the role to a specific database and avoid broad CREATEDB permissions after deployment.

Grants are scoped to databases, schemas, tables, sequences, and functions. Granting a database privilege does not automatically give table privileges inside every schema.

DB02 · TEACHING 3

psql CLI, Users, Roles, and Grants (continued)

Password authentication requires both a password on the role and a `pg_hba.conf` rule that permits password-based access. When either side is missing, Odoo errors often look like a generic authentication failure.

Use psql meta commands such as `\du`, `\l`, `\dn`, and `\dt` to inspect the server quickly. These commands are not SQL, but they are invaluable during support and training.

DB02 · KEY POINTS

What to remember

- A PostgreSQL role is the central identity and privilege unit.
- LOGIN makes a role usable as a user account.
- CREATEDB is convenient for Odoo development but should be reviewed for production.
- psql meta commands are part of a practical database debugging workflow.

DB02 · EXAMPLE 1

Create an Odoo database role

Create an Odoo database role

This creates a login role that can create databases, then confirms that password authentication works over TCP.

```
sudo -u postgres createuser --createdb --pwprompt odoo17
psql -h 127.0.0.1 -U odoo17 -d postgres

-- Inside psql
\conninfo
\du
SELECT current_user;
```

DB02 · EXAMPLE 2

Grant schema privileges

Grant schema privileges

The example separates ownership from application access, which is closer to how production systems are often controlled.

```
CREATE DATABASE training_owner OWNER postgres;
CREATE ROLE training_app LOGIN PASSWORD 'change-me';
\c training_owner
GRANT USAGE, CREATE ON SCHEMA public TO training_app;
ALTER DEFAULT PRIVILEGES IN SCHEMA public
    GRANT SELECT, INSERT, UPDATE, DELETE ON TABLES TO training_app;
```

DB02 · PITFALLS

Common mistakes

- Assuming a database grant automatically grants privileges on every table.
- Forgetting that peer authentication uses the Linux user name, not the typed database password.
- Giving SUPERUSER to application roles to avoid learning the real permission issue.
- Leaving training passwords in shell history or shared notes.

Remember: Validate fixes in a disposable database before production.

DB03 · OBJECTIVES

Create and Manage Databases

- Create databases with the correct owner and encoding.
- Understand templates, connection limits, and database ownership.
- Prepare databases for Odoo import and development workflows.

Lab target: Create two training databases owned by your Odoo role, list them with \l, and remove one after confirming its owner.

DB03 · TEACHING 1

Create and Manage Databases

A PostgreSQL database is an isolation boundary for Odoo data. Each Odoo database contains its own registry metadata, business records, installed modules, and attachments references.

The database owner matters because Odoo creates and updates many database objects during module installation and upgrades. If the role lacks ownership or enough privileges, upgrade failures can appear deep inside ORM operations.

DB03 · TEACHING 2

Create and Manage Databases (continued)

UTF8 encoding is the normal choice for Odoo. Training environments should make this explicit so later collation, translation, and text search behavior are easier to explain.

Template databases allow PostgreSQL to copy an existing database structure at creation time. For ordinary Odoo work, `template0` or the default `template1` is enough, but polluted templates can introduce surprising extensions or objects.

DB03 · TEACHING 3

Create and Manage Databases (continued)

Connection limits can protect a shared database server from one workload consuming every backend process. For local training, low limits are annoying, but for production they can prevent cascading failure.

Dropping an Odoo database is destructive and should be treated like deleting a customer environment. Always confirm you have the right database name and a tested backup before destructive maintenance.

DB03 · KEY POINTS

What to remember

- One Odoo database maps to one Odoo instance dataset.
- Use explicit owners for predictable module installation and upgrades.
- UTF8 is expected for modern Odoo deployments.
- Do not drop or rename databases casually in shared environments.

DB03 · EXAMPLE 1

Create databases from SQL and CLI

Create databases from SQL and CLI

The CLI is concise for daily work, while SQL makes ownership, encoding, template, and limits explicit.

```
createdb -h 127.0.0.1 -U odoo17 --encoding=UTF8 odoo17_training
```

```
psql -U postgres -d postgres -c "  
CREATE DATABASE odoo17_reporting  
    WITH OWNER odoo17  
    ENCODING 'UTF8'  
    TEMPLATE template0  
    CONNECTION LIMIT 50;  
"
```

DB03 · PITFALLS

Common mistakes

- Creating an Odoo database owned by postgres and then running Odoo as a different restricted role.
- Using mixed encodings or locale settings across environments without documenting them.
- Dropping a database from the wrong terminal session.
- Copying a database while users or cron jobs are still writing to it.

Remember: Validate fixes in a disposable database before production.

DB03 · EXTRA

Odoo habit

Odoo habit: Name development databases clearly, for example `odoo17_dev_customer_a`, so backups and logs are easier to match.

DB04 · OBJECTIVES

PostgreSQL Data Types for Odoo Developers

- Recognize PostgreSQL types commonly produced by Odoo fields.
- Choose SQL types for reporting tables and integration staging tables.
- Avoid type choices that damage money, date, or JSON accuracy.

Lab target: Inspect five columns from res_partner or another Odoo table and map each SQL type back to the likely Odoo field type.

DB04 · TEACHING 1

PostgreSQL Data Types for Odoo Developers

Odoo hides most column definitions behind its ORM, but PostgreSQL types still affect performance, correctness, and reporting. Senior Odoo developers should be able to inspect a table and understand what the ORM generated.

Text-like fields commonly map to varchar or text. The practical distinction is less about storage and more about validation, indexing, and how the Odoo field communicates intent to users.

Numeric choices matter for money and quantities. Use numeric for exact decimal values when precision matters, and avoid float for values that must reconcile financially.

DB04 · TEACHING 2

PostgreSQL Data Types for Odoo Developers (continued)

Date and timestamp columns support business workflows such as accounting periods, scheduled actions, and SLA calculations. Odoo stores datetimes in UTC and converts them for users at the application layer.

JSONB is useful for flexible integration payloads, analytic properties, and denormalized metadata. It is powerful, but it should not become a hiding place for core relational data that needs constraints and joins.

DB04 · TEACHING 3

PostgreSQL Data Types for Odoo Developers (continued)

Arrays, UUIDs, and specialized types can be useful in custom SQL tables, but they should be introduced deliberately. Odoo's ORM works best when custom models use standard field types and clear relationships.

DB04 · KEY POINTS

What to remember

- Understand the SQL type behind each important Odoo field.
- Use numeric for exact business decimals.
- Datetimes are stored in UTC by Odoo.
- JSONB is useful, but relational data still belongs in relations.

DB04 · EXAMPLE 1

Inspect generated Odoo column types

Inspect generated Odoo column types

The query teaches developers to inspect real Odoo tables, while the staging table shows conservative type choices for integrations.

```
SELECT column_name, data_type, is_nullable
FROM information_schema.columns
WHERE table_name = 'res_partner'
ORDER BY ordinal_position;

CREATE TABLE integration_payload (
  id bigserial PRIMARY KEY,
  external_id varchar NOT NULL,
  amount numeric(16, 2) NOT NULL,
  received_at timestamp without time zone NOT NULL,
  payload jsonb NOT NULL
);
```

DB04 · PITFALLS

Common mistakes

- Using float for money because it looks simpler in test data.
- Storing dates as text and then losing reliable filtering and sorting.
- Putting frequently joined business identifiers only inside JSONB.
- Assuming Odoo user timezone conversion changes how datetimes are stored in PostgreSQL.

Remember: Validate fixes in a disposable database before production.

DB05 · OBJECTIVES

DDL: CREATE, ALTER, and DROP

- Use DDL safely in development and maintenance tasks.
- Understand how DDL locks affect live Odoo systems.
- Plan schema changes with reversibility and backups in mind.

Lab target: Create a reporting schema and table, alter it with one new column, inspect it with \d, then drop it after confirming the object name.

DB05 · TEACHING 1

DDL: CREATE, ALTER, and DROP

Data Definition Language changes the structure of the database. In Odoo, most application schema changes should come from module code and migrations, not from manual SQL typed into production.

CREATE statements add tables, indexes, schemas, views, and functions. They are safe when planned, but object names should follow clear conventions so DBAs and developers can identify ownership.

DB05 · TEACHING 2

DDL: CREATE, ALTER, and DROP (continued)

ALTER statements can be cheap or expensive depending on the operation and PostgreSQL version. Adding a nullable column is usually fast, while changing a type or setting a not-null constraint may scan or rewrite data.

DROP statements are final from the perspective of the current transaction once committed. A backup is not optional when dropping objects that may contain customer or audit data.

DDL takes locks, and those locks can block Odoo workers. Even short DDL can create visible downtime if it waits behind a long transaction and then blocks new writes.

DB05 · TEACHING 3

DDL: CREATE, ALTER, and DROP (continued)

For custom Odoo modules, prefer declarative model fields and migration scripts that can be reviewed. Manual DDL is best reserved for diagnostics, reporting objects, and carefully controlled maintenance.

DB05 · KEY POINTS

What to remember

- Let Odoo module definitions manage normal model tables.
- DDL can block live application traffic.
- Review ALTER operations for rewrite or scan behavior.
- Treat DROP as a backup-required operation.

DB05 · EXAMPLE 1

DDL for a reporting helper table

DDL for a reporting helper table

This illustrates common DDL statements in a non-Odoo schema where manual SQL may be appropriate.

```
CREATE SCHEMA IF NOT EXISTS reporting;

CREATE TABLE reporting.daily_sales_snapshot (
    snapshot_date date PRIMARY KEY,
    order_count integer NOT NULL DEFAULT 0,
    untaxed_total numeric(16, 2) NOT NULL DEFAULT 0
);

ALTER TABLE reporting.daily_sales_snapshot
    ADD COLUMN company_id integer;

DROP TABLE reporting.daily_sales_snapshot;
```

DB05 · PITFALLS

Common mistakes

- Manually altering Odoo-owned tables and then being surprised when module upgrades disagree.
- Running ALTER TABLE during business hours without checking locks and table size.
- Dropping a column before confirming whether computed fields, reports, or integrations use it.
- Leaving experimental tables in public schema with unclear ownership.

Remember: Validate fixes in a disposable database before production.

LESSONS DB06–DB10

PostgreSQL Administration & SQL

Teaching block 2

DB06 · OBJECTIVES

Primary Keys, Foreign Keys, and Constraints

- Explain how constraints protect relational integrity.
- Use primary keys, foreign keys, unique constraints, and checks.
- Connect database constraints to Odoo model design.

Lab target: Build two related tables with a foreign key and a unique constraint, then attempt inserts that prove each constraint works.

DB06 · TEACHING 1

Primary Keys, Foreign Keys, and Constraints

Constraints are the database's way of refusing impossible data. They remain active regardless of whether data is written by Odoo, an import script, psql, or an integration job.

A primary key uniquely identifies each row and is the target for many relationships. Odoo models normally use an integer id primary key generated by the database.

Foreign keys preserve references between tables. In Odoo, Many2one fields often create foreign keys unless the field definition or module history prevents it.

DB06 · TEACHING 2

Primary Keys, Foreign Keys, and Constraints (continued)

Unique constraints prevent duplicates that application checks can miss under concurrency. For example, two users creating the same external reference at the same time should be stopped by the database, not only by Python code.

Check constraints enforce local row rules such as positive quantities or valid ranges. They are simple, fast, and useful for invariants that should never be bypassed.

DB06 · TEACHING 3

Primary Keys, Foreign Keys, and Constraints (continued)

Good Odoo modules combine ORM validations with SQL constraints. Python constraints give friendly messages and cross-record logic, while SQL constraints provide final protection against race conditions.

DB06 · KEY POINTS

What to remember

- Constraints are stronger than application-only validation.
- Unique constraints are important under concurrent writes.
- Foreign keys keep relational references honest.
- Use SQL constraints for invariants that must never be violated.

DB06 · EXAMPLE 1

Relational constraints

Relational constraints

The schema protects identity, references, positive amounts, and invoice uniqueness even if data is loaded outside Odoo.

```
CREATE TABLE training_customer (  
    id bigserial PRIMARY KEY,  
    email varchar NOT NULL UNIQUE,  
    credit_limit numeric(16, 2) NOT NULL DEFAULT 0,  
    CHECK (credit_limit >= 0)  
);  
  
CREATE TABLE training_invoice (  
    id bigserial PRIMARY KEY,  
    customer_id bigint NOT NULL REFERENCES training_customer(id) ON DELETE RESTRICT,  
    number varchar NOT NULL,  
    total numeric(16, 2) NOT NULL CHECK (total >= 0),  
    UNIQUE (customer_id, number)  
);
```

DB06 · PITFALLS

Common mistakes

- Relying only on Python validation for uniqueness in concurrent workflows.
- Using ON DELETE CASCADE without considering accounting and audit requirements.
- Adding constraints to dirty data without first finding violations.
- Creating nullable foreign keys when the business relationship is actually required.

Remember: Validate fixes in a disposable database before production.

DB07 · OBJECTIVES

CRUD: INSERT, SELECT, UPDATE, and DELETE

- Write basic CRUD SQL confidently.
- Understand transactions and RETURNING clauses.
- Relate direct SQL changes to Odoo ORM safety rules.

Lab target: Insert a row in a transaction, update it with RETURNING, verify the result, and roll back so the table returns to its original state.

DB07 · TEACHING 1

CRUD: INSERT, SELECT, UPDATE, and DELETE

CRUD operations are the foundation of every database-backed application. Even when Odoo's ORM writes SQL for you, understanding the underlying statements helps you debug logs and performance plans.

INSERT creates rows and can return generated values such as id. RETURNING is especially useful in scripts because it avoids a second query to discover what was inserted.

SELECT retrieves data and should be shaped intentionally. Pulling every column from a large Odoo table is rarely the right diagnostic habit when you only need ids, names, and a status field.

DB07 · TEACHING 2

CRUD: INSERT, SELECT, UPDATE, and DELETE (continued)

UPDATE changes existing rows and is dangerous without a precise WHERE clause. In Odoo databases, direct UPDATE can bypass computed fields, mail tracking, access rules, onchange logic, and business validations.

DELETE removes rows and should be used with even more caution. Many Odoo records should be archived with `active = false` rather than deleted, especially if they are referenced by accounting, inventory, or chatter history.

DB07 · TEACHING 3

CRUD: INSERT, SELECT, UPDATE, and DELETE (continued)

Transactions let several statements succeed or fail together. In psql, BEGIN, COMMIT, and ROLLBACK are essential safety tools when testing data changes.

DB07 · KEY POINTS

What to remember

- Always write SELECT first to confirm the target rows before UPDATE or DELETE.
- RETURNING is useful for inserts and updates.
- Direct SQL bypasses Odoo ORM behavior.
- Use transactions for safe experimentation.

DB07 · EXAMPLE 1

Safe CRUD workflow

Safe CRUD workflow

The transaction proves the statements and then rolls them back, which is a safe habit when learning or diagnosing.

```
BEGIN;

INSERT INTO training_customer (email, credit_limit)
VALUES ('ada@example.com', 5000.00)
RETURNING id, email;

SELECT id, email, credit_limit
FROM training_customer
WHERE email = 'ada@example.com';

UPDATE training_customer
SET credit_limit = credit_limit + 1000
WHERE email = 'ada@example.com'
RETURNING id, credit_limit;

ROLLBACK;
```


DB07 · PITFALLS

Common mistakes

- Running UPDATE or DELETE without a WHERE clause.
- Changing Odoo data directly and expecting ORM computed fields or tracking to run.
- Using SELECT * on large tables during production incidents.
- Forgetting to COMMIT a successful transaction or ROLLBACK a test transaction.

Remember: Validate fixes in a disposable database before production.

DB08 · OBJECTIVES

Filtering with WHERE, LIKE, ILIKE, IN, and BETWEEN

- Use common filtering operators precisely.
- Understand case-sensitive and case-insensitive pattern matching.
- Avoid filters that defeat indexes or return unintended records.

Lab target: Write three SELECT queries against an Odoo database: one with ILIKE, one with IN, and one date range using >= and <.

DB08 · TEACHING 1

Filtering with WHERE, LIKE, ILIKE, IN, and BETWEEN

Filtering is where SQL becomes practical for support work. Good WHERE clauses narrow the problem to the exact company, date range, state, or external identifier you need to investigate.

LIKE is case-sensitive pattern matching in PostgreSQL, while ILIKE is case-insensitive. ILIKE is convenient for support searches but can be slower unless supported by the right index strategy.

IN is clearer than a long chain of OR comparisons when matching a known set of values. It is common in Odoo diagnostics because record ids are often copied from logs, URLs, or failed queue jobs.

DB08 · TEACHING 2

Filtering with WHERE, LIKE, ILIKE, IN, and BETWEEN (continued)

BETWEEN includes both boundaries. That is useful for dates when you mean a closed range, but for datetimes many teams prefer `>= start` and `< next_start` to avoid edge mistakes.

NULL requires `IS NULL` or `IS NOT NULL`, not equals comparisons. This distinction matters in Odoo because optional Many2one fields are stored as NULL when unset.

Filters should be written with indexes in mind when tables are large. Wrapping indexed columns in functions or starting text patterns with a wildcard can force PostgreSQL to inspect far more rows.

DB08 · KEY POINTS

What to remember

- LIKE is case-sensitive; ILIKE is case-insensitive.
- BETWEEN includes both ends of the range.
- Use IS NULL for missing optional values.
- Index-friendly filters matter on large Odoo tables.

DB08 · EXAMPLE 1

Odoo-style support filters

Odoo-style support filters

The examples use practical filters for partners and orders while avoiding an ambiguous datetime BETWEEN boundary.

```
SELECT id, name, email
FROM res_partner
WHERE active IS TRUE
      AND email ILIKE '%@example.com'
      AND company_id IN (1, 3, 7)
ORDER BY name
LIMIT 50;

SELECT id, name, date_order
FROM sale_order
WHERE date_order >= timestamp '2026-01-01'
      AND date_order < timestamp '2026-02-01'
      AND state IN ('sale', 'done');
```

DB08 · PITFALLS

Common mistakes

- Expecting LIKE to be case-insensitive.
- Using BETWEEN with datetimes and accidentally including or excluding boundary records.
- Writing column = NULL instead of column IS NULL.
- Starting every text search with % and then wondering why indexes are not used.

Remember: Validate fixes in a disposable database before production.

DB09 · OBJECTIVES

Joins for Relational Odoo Data

- Choose INNER JOIN and LEFT JOIN correctly.
- Trace Odoo Many2one and Many2many relationships in SQL.
- Read joined results without multiplying rows accidentally.

Lab target: Join sale_order to res_partner, then write a separate query that follows one Many2many relation table in your database.

DB09 · TEACHING 1

Joins for Relational Odoo Data

Joins combine rows from related tables, which is essential in Odoo because business data is highly relational. A sale order links to partners, companies, currencies, users, order lines, stock moves, invoices, and more.

INNER JOIN returns rows only when both sides match. It is appropriate when the relationship is required and missing related data would indicate invalid or irrelevant records.

LEFT JOIN keeps the left-side row even when the related row is missing. This is useful for optional fields, incomplete setup, and diagnostics where missing data is part of the question.

DB09 · TEACHING 2

Joins for Relational Odoo Data (continued)

Many2one relationships are usually simple joins from a foreign key column to the related table's id. Many2many relationships use an intermediate relation table that contains pairs of ids.

One-to-many joins can multiply rows because each parent may have many children. That is expected, but it becomes a bug when totals are summed after joining to another one-to-many table without pre-aggregation.

When debugging Odoo reports, inspect the join path carefully before blaming the ORM. Many incorrect totals come from SQL joins that accidentally duplicate business rows.

DB09 · KEY POINTS

What to remember

- INNER JOIN filters to matching rows.
- LEFT JOIN preserves unmatched left-side rows.
- Many2many joins require a relation table.
- One-to-many joins can duplicate parent values in result sets.

DB09 · EXAMPLE 1

Join sale orders to customers and tags

Join sale orders to customers and tags

The first query follows Many2one fields, and the second shows the relation table pattern behind a partner tag Many2many.

```
SELECT so.id, so.name, rp.name AS customer, rc.name AS company
FROM sale_order so
JOIN res_partner rp ON rp.id = so.partner_id
LEFT JOIN res_company rc ON rc.id = so.company_id
WHERE so.state IN ('sale', 'done')
ORDER BY so.date_order DESC
LIMIT 20;
```

```
SELECT rp.name, category.name AS tag
FROM res_partner rp
JOIN res_partner_res_partner_category_rel rel
  ON rel.partner_id = rp.id
JOIN res_partner_category category
  ON category.id = rel.category_id
WHERE rp.id = 42;
```

DB09 · PITFALLS

Common mistakes

- Using INNER JOIN where optional data requires LEFT JOIN.
- Summing parent totals after joining multiple child tables.
- Guessing Many2many relation table names instead of inspecting the model or database.
- Joining on display names instead of stable ids.

Remember: Validate fixes in a disposable database before production.

DB10 · OBJECTIVES

GROUP BY, HAVING, and Aggregates

- Summarize Odoo data with aggregate functions.
- Use GROUP BY and HAVING in the correct order.
- Avoid aggregation mistakes caused by joins and NULL values.

Lab target: Create a grouped sales or invoice query with COUNT and SUM, then add a HAVING clause that filters low-value groups.

DB10 · TEACHING 1

GROUP BY, HAVING, and Aggregates

Aggregates turn detail records into business signals such as totals, counts, minimums, maximums, and averages. Odoo dashboards and reports often reduce to the same concepts expressed through ORM `read_group` or SQL.

GROUP BY defines the grain of the result. If you group sales by company and month, every selected non-aggregate column must belong to that grain.

WHERE filters rows before grouping, while HAVING filters groups after aggregation. This difference is important when you want only customers whose total confirmed sales exceed a threshold.

DB10 · TEACHING 2

GROUP BY, HAVING, and Aggregates (continued)

COUNT(*) counts rows, while COUNT(column) counts non-null values. In optional Odoo fields, the difference can change your conclusion during data quality checks.

Aggregating after joins requires care. If a sale order is joined to both order lines and invoices, totals may be multiplied unless one side is aggregated first.

For application code, prefer Odoo's `read_group` when you need ORM security, context, and computed grouping behavior. Use SQL for diagnostics, migrations, and carefully reviewed reporting queries.

DB10 · KEY POINTS

What to remember

- GROUP BY sets the result grain.
- HAVING filters aggregate results.
- COUNT(*) and COUNT(column) answer different questions.
- Pre-aggregate child tables before joining multiple one-to-many relationships.

DB10 · EXAMPLE 1

Monthly sales summary

Monthly sales summary

The query filters confirmed orders first, groups by month and company, then keeps only meaningful sales groups.

```
SELECT
    date_trunc('month', so.date_order)::date AS month,
    so.company_id,
    COUNT(*) AS order_count,
    SUM(so.amount_total) AS total_amount
FROM sale_order so
WHERE so.state IN ('sale', 'done')
GROUP BY 1, 2
HAVING SUM(so.amount_total) > 10000
ORDER BY month DESC, total_amount DESC;
```

DB10 · PITFALLS

Common mistakes

- Selecting columns that are not part of GROUP BY and not aggregated.
- Using HAVING for simple row filters that belong in WHERE.
- Forgetting that COUNT(column) ignores NULL.
- Joining detail tables before aggregation and inflating totals.

Remember: Validate fixes in a disposable database before production.

DB10 · EXTRA

Aggregate order

Clause	Purpose
WHERE	Filter detail rows before grouping
GROUP BY	Define aggregate grain
HAVING	Filter aggregate groups
ORDER BY	Sort the final result

LESSONS DB11–DB15

PostgreSQL Administration & SQL

Teaching block 3

DB11 · OBJECTIVES

Indexes: B-Tree, GIN, and Composite Design

- Explain how indexes speed up reads and slow down writes.
- Choose B-Tree, GIN, and composite indexes for common cases.
- Design indexes around Odoo domains and reporting queries.

Lab target: Choose one slow-looking domain from Odoo logs, write the SQL equivalent, run EXPLAIN ANALYZE, add an index in a test database, and compare the plan.

DB11 · TEACHING 1

Indexes: B-Tree, GIN, and Composite Design

Indexes are lookup structures that help PostgreSQL find rows without scanning an entire table. They are essential for large Odoo installations because tables such as `account_move_line`, `mail_message`, and `stock_move` can grow quickly.

B-Tree indexes are the default and support equality, range comparisons, ordering, and many common joins. They are usually the right first choice for ids, foreign keys, dates, states, and unique constraints.

DB11 · TEACHING 2

Indexes: B-Tree, GIN, and Composite Design (continued)

Composite indexes include more than one column and should match real filter patterns. An index on (company_id, state, date_order) can help a domain that filters company and state before sorting or ranging by date.

GIN indexes are useful for containment and full-text style lookups, especially with jsonb, arrays, and trigram search when the extension is enabled. They are heavier than B-Tree indexes, so use them when the query pattern justifies the cost.

DB11 · TEACHING 3

Indexes: B-Tree, GIN, and Composite Design (continued)

Every index has a write cost because PostgreSQL must maintain it during INSERT, UPDATE, and DELETE. Over-indexing an Odoo table can slow imports, stock validation, accounting posting, and module upgrades.

Good index work starts from a slow query or a known high-volume domain, not from guesswork. Use EXPLAIN ANALYZE and production-like data before adding indexes that will live for years.

DB11 · KEY POINTS

What to remember

- B-Tree is the default index type for most relational filters.
- Composite index column order should match query patterns.
- GIN is useful for jsonb, arrays, and specialized text search.
- Indexes speed reads but add write and storage cost.

DB11 · EXAMPLE 1

Indexes for Odoo query patterns

Indexes for Odoo query patterns

The examples show B-Tree indexes for common domains and a GIN index for JSONB containment queries.

```
CREATE INDEX sale_order_company_state_date_idx  
ON sale_order (company_id, state, date_order DESC);
```

```
CREATE INDEX account_move_line_partner_date_idx  
ON account_move_line (partner_id, date);
```

```
CREATE INDEX integration_payload_payload_gin_idx  
ON integration_payload USING gin (payload);
```

```
EXPLAIN ANALYZE  
SELECT id, name  
FROM sale_order  
WHERE company_id = 1  
      AND state IN ('sale', 'done')  
ORDER BY date_order DESC  
LIMIT 50;
```

DB11 · PITFALLS

Common mistakes

- Adding indexes without measuring the query before and after.
- Creating many overlapping indexes on busy Odoo tables.
- Putting composite index columns in an order that does not match real filters.
- Expecting a normal B-Tree index to make leading-wildcard ILIKE searches fast.

Remember: Validate fixes in a disposable database before production.

DB12 · OBJECTIVES

Reading EXPLAIN ANALYZE

- Use EXPLAIN ANALYZE to inspect real query execution.
- Interpret scans, joins, row estimates, timing, and buffers.
- Turn query plans into practical optimization decisions.

Lab target: Run EXPLAIN (ANALYZE, BUFFERS) for a filtered query before and after ANALYZE, then document the scan type and row estimates.

DB12 · TEACHING 1

Reading EXPLAIN ANALYZE

EXPLAIN shows PostgreSQL's planned execution strategy, while EXPLAIN ANALYZE runs the query and reports what actually happened. The ANALYZE form is more truthful but must be used carefully for writes because it executes them.

Plans are trees, and each node contributes work to its parent. Sequential scans, index scans, bitmap scans, nested loops, hash joins, and sorts all tell you how the optimizer chose to access and combine data.

DB12 · TEACHING 2

Reading EXPLAIN ANALYZE (continued)

Estimated rows versus actual rows is one of the most important signals. Large differences can mean stale statistics, correlated filters, missing indexes, or a query shape the planner cannot estimate well.

Timing tells you where time was spent, but buffers often explain why. A plan that reads many shared buffers may be doing heavy logical I/O even if the current run is warmed by cache.

An index scan is not automatically good, and a sequential scan is not automatically bad. Scanning a small table or a large percentage of a table can be cheaper than bouncing through an index.

DB12 · TEACHING 3

Reading EXPLAIN ANALYZE (continued)

For Odoo, connect slow SQL logs to EXPLAIN work. The best optimizations usually come from understanding a repeated ORM domain, adding a targeted index, reducing selected columns, or avoiding an accidental join explosion.

DB12 · KEY POINTS

What to remember

- EXPLAIN ANALYZE executes the query.
- Compare estimated rows to actual rows.
- Look at buffers as well as timing.
- Sequential scans can be correct for small or broad queries.

DB12 · EXAMPLE 1

Analyze a read query safely

Analyze a read query safely

The first command reports actual execution details, and the second refreshes statistics if estimates look stale.

```
EXPLAIN (ANALYZE, BUFFERS, VERBOSE)
SELECT id, name, date_order
FROM sale_order
WHERE company_id = 1
      AND state = 'sale'
ORDER BY date_order DESC
LIMIT 20;

ANALYZE sale_order;
```

DB12 · PITFALLS

Common mistakes

- Running EXPLAIN ANALYZE on UPDATE or DELETE in production without wrapping and rolling back.
- Assuming the largest cost number is always the real bottleneck.
- Ignoring row estimate errors.
- Adding an index because a sequential scan appears, even when the query returns most rows.

Remember: Validate fixes in a disposable database before production.

DB12 · EXTRA

Plan reading checklist

- Which node reads the most rows?
- Where do estimates differ from actual rows?
- Is the query sorting, hashing, or looping more than expected?
- Are buffers high for a query that should be selective?

DB13 · OBJECTIVES

Backups with pg_dump

- Create logical backups with pg_dump.
- Choose custom, directory, and plain SQL dump formats.
- Prepare Odoo databases for consistent backup routines.

Lab target: Create a custom-format dump of a training database and record the dump size, duration, and exact command used.

DB13 · TEACHING 1

Backups with pg_dump

pg_dump creates a logical backup by reading database objects and data. It is portable across machines and useful for Odoo migrations, staging refreshes, and point-in-time training snapshots.

The custom format is usually the best default for Odoo because it is compressed, flexible, and restorable with pg_restore. Plain SQL is readable but less flexible for selective restore and parallel jobs.

A consistent backup must include the PostgreSQL database and Odoo filestore. The database stores attachment metadata, but the binary files may live on disk under the Odoo data directory.

DB13 · TEACHING 2

Backups with pg_dump (continued)

Backups should be automated and verified, not performed only when someone remembers. A backup that has never been restored is an assumption, not a recovery plan.

Large databases need attention to runtime, disk space, compression, and impact on the database server. pg_dump is online, but it still consumes I/O and CPU that can affect users.

Security matters because dumps contain customer data, passwords hashes, tokens, messages, and accounting records. Encrypt backups at rest and control who can download them.

DB13 · KEY POINTS

What to remember

- Use custom format for flexible pg_restore workflows.
- Back up the Odoo filestore with the database.
- Automate backups and test restores.
- Protect dumps as sensitive production data.

DB13 · EXAMPLE 1

Custom-format Odoo backup

Custom-format Odoo backup

The database dump and filestore archive are both needed for a realistic Odoo recovery.

```
export DB=odoo17_prod
export BACKUP_DIR=/var/backups/odoo
sudo mkdir -p "$BACKUP_DIR"

pg_dump -h 127.0.0.1 -U odoo17 -F c -Z 6      -f "$BACKUP_DIR/${DB}_${date +%F}.dump" "$DB"

sudo tar -C /var/lib/odoo/.local/share/Odoo/filestore      -czf "$BACKUP_DIR/${DB}_filestore_${date +%F}.tar.gz" "$DB"
```

DB13 · PITFALLS

Common mistakes

- Backing up only PostgreSQL and forgetting the Odoo filestore.
- Using plain SQL dumps for every case and losing pg_restore flexibility.
- Keeping backups on the same disk as the database with no off-host copy.
- Never testing whether the backup can be restored.

Remember: Validate fixes in a disposable database before production.

DB14 · OBJECTIVES

Restores with pg_restore

- Restore custom-format dumps safely.
- Use clean, create, jobs, and selective restore options.
- Validate an Odoo database after restore.

Lab target: Restore a dump into a new training database, start Odoo against it, and open one attachment to prove the filestore is present.

DB14 · TEACHING 1

Restores with pg_restore

Restore practice is where backup confidence becomes real. A team that restores regularly can estimate recovery time, detect missing filestore data, and find permission problems before an emergency.

pg_restore works with custom and directory-format dumps. It can create objects, clean existing objects, run parallel jobs, and restore selected schemas or tables when the dump format supports it.

DB14 · TEACHING 2

Restores with pg_restore (continued)

Restoring into a fresh database is usually safer than overwriting an existing one. It lets you validate the result before swapping names, updating configuration, or pointing Odoo at the restored database.

Ownership and role availability matter during restore. If the dump references owners that do not exist on the target server, you may need `--no-owner` or a prepared role setup.

For Odoo, restore the filestore directory with the exact database name expected by the instance. If the database is renamed, the filestore path must match or attachments will appear missing.

DB14 · TEACHING 3

Restores with pg_restore (continued)

After restore, validate at multiple levels: psql connection, Odoo registry startup, module list, attachment access, scheduled actions, and a few business workflows. A database that restores without SQL errors may still be incomplete operationally.

DB14 · KEY POINTS

What to remember

- Practice restore before a real incident.
- Prefer restoring into a fresh database for validation.
- Use --no-owner when target roles differ.
- Match the filestore directory name to the restored Odoo database.

DB14 · EXAMPLE 1

Restore a custom dump to a test database

Restore a custom dump to a test database

The restore targets a separate database and performs a simple Odoo metadata check afterward.

```
createdb -h 127.0.0.1 -U odoo17 odoo17_restore_test

pg_restore -h 127.0.0.1 -U odoo17      --dbname=odoo17_restore_test      --no-owner      --jobs=4      /var/backups/odoo/odoo17_pro

psql -h 127.0.0.1 -U odoo17 -d odoo17_restore_test      -c "SELECT latest_version FROM ir_module_module WHERE name = 'base';"
```


DB14 · PITFALLS

Common mistakes

- Restoring over the only available copy of a database.
- Ignoring missing target roles and ownership errors.
- Forgetting to restore or rename the filestore.
- Declaring success after `pg_restore` exits without checking Odoo startup.

Remember: Validate fixes in a disposable database before production.

DB15 · OBJECTIVES

postgresql.conf and pg_hba.conf

- Locate and edit the main PostgreSQL configuration files.
- Configure listening addresses, memory settings, logging, and authentication.
- Reload or restart PostgreSQL safely after configuration changes.

Lab target: Find config_file and hba_file with SHOW commands, add a localhost rule for your Odoo role in a lab environment, reload, and reconnect.

DB15 · TEACHING 1

postgresql.conf and pg_hba.conf

postgresql.conf controls server behavior such as listening addresses, memory, logging, autovacuum, and connection limits.

pg_hba.conf controls who can connect, from where, to which database, and by which authentication method.

Configuration files are cluster-specific on Ubuntu. The common path includes the PostgreSQL major version, such as /etc/postgresql/15/main, so always confirm the active cluster before editing.

DB15 · TEACHING 2

postgresql.conf and pg_hba.conf (continued)

listen_addresses determines the network interfaces PostgreSQL listens on. Binding to localhost is safest for single-host Odoo; binding to all addresses requires firewalling and strong authentication rules.

pg_hba.conf is evaluated top to bottom, and the first matching rule wins. A permissive rule above a strict rule can silently weaken the configuration.

Many changes only need a reload, while others require a restart. Reloading is less disruptive, but you must know whether the setting is reloadable before assuming it took effect.

DB15 · TEACHING 3

postgresql.conf and pg_hba.conf (continued)

Odoo deployments benefit from logging slow queries and connection information during tuning. Do not enable extremely verbose logging permanently without checking disk growth and privacy implications.

DB15 · KEY POINTS

What to remember

- postgresql.conf tunes server behavior.
- pg_hba.conf controls connection authentication.
- pg_hba.conf rule order matters.
- Reload when possible, restart when required.

DB15 · EXAMPLE 1

Odoo-friendly local configuration snippets

Odoo-friendly local configuration snippets

The snippets keep PostgreSQL local to the application host and enable slow-query visibility for tuning.

```
-- Locate files
SHOW config_file;
SHOW hba_file;

-- postgresql.conf examples
listen_addresses = '127.0.0.1'
max_connections = 200
shared_buffers = 1GB
log_min_duration_statement = 500

-- pg_hba.conf example
# TYPE      DATABASE      USER      ADDRESS      METHOD
host       all             odoo17     127.0.0.1/32  scram-sha-256

-- Apply reloadable changes
SELECT pg_reload_conf();
```

DB15 · PITFALLS

Common mistakes

- Editing the wrong versioned cluster directory.
- Adding a strict `pg_hba.conf` rule below a broader rule that already matches.
- Using trust authentication on shared or production machines.
- Restarting the service for every change without knowing whether reload is enough.

Remember: Validate fixes in a disposable database before production.

LESSONS DB16–DB16

PostgreSQL Administration & SQL

Teaching block 4

DB16 · OBJECTIVES

Locks, Connections, Terminating Queries, and Odoo Backup Habits

- Inspect active connections and blocking locks.
- Terminate unsafe or stuck sessions carefully.
- Apply Odoo-specific database habits and backup strategy.

Lab target: Open two psql sessions, create a blocking transaction in one, observe it from pg_stat_activity in the other, then cancel or roll back safely.

DB16 · TEACHING 1

Locks, Connections, Terminating Queries, and Odoo Backup Habits

Every PostgreSQL connection uses server resources, and every transaction can hold locks. Odoo workers, cron jobs, imports, reports, and shell sessions can all contribute to lock contention.

Locks are normal; blocking is the operational concern. A short row lock during a write is expected, while an idle transaction holding a table lock can stop module upgrades or business workflows.

`pg_stat_activity` shows what sessions are doing, how long they have been running, and whether they are waiting. Combined with `pg_locks`, it helps identify blockers rather than only victims.

DB16 · TEACHING 2

Locks, Connections, Terminating Queries, and Odoo Backup Habits (continued)

Terminating a backend is a production action, not a reflex. Prefer canceling a query when possible, communicate with users when needed, and understand whether the application will retry safely.

Odoo-specific database habits include avoiding direct writes to business tables, stopping Odoo during risky restores, disabling cron during staging refreshes, and running upgrades from the application layer. Respect the ORM unless you are deliberately doing maintenance.

DB16 · TEACHING 3

Locks, Connections, Terminating Queries, and Odoo Backup Habits (continued)

A strong Odoo backup strategy combines scheduled `pg_dump` or physical backups, filestore backups, off-host copies, retention, encryption, and restore drills. The strategy should define recovery point objective, recovery time objective, and who is allowed to perform recovery.

DB16 · KEY POINTS

What to remember

- Use `pg_stat_activity` to understand sessions before acting.
- Find blockers, not just blocked queries.
- Cancel or terminate with intent and an audit trail.
- Odoo backups must include both database and filestore.
- Restore drills are part of the backup strategy.

DB16 · EXAMPLE 1

Inspect and control active sessions

Inspect and control active sessions

Inspect first, cancel when possible, and terminate only when the risk of leaving the session running is greater than killing it.

```
SELECT pid, username, datname, state, wait_event_type, wait_event,  
       now() - query_start AS runtime,  
       left(query, 120) AS query  
FROM pg_stat_activity  
WHERE datname = 'odoo17_prod'  
ORDER BY runtime DESC NULLS LAST;  
  
SELECT pg_cancel_backend(12345);  
SELECT pg_terminate_backend(12345);
```

DB16 · EXAMPLE 2

Simple Odoo backup policy outline

Simple Odoo backup policy outline

The outline connects backup creation to restore testing, which is the habit that actually proves recoverability.

Daily:

- `pg_dump` custom format for each production database
- tar and compress matching filestore directories
- upload encrypted artifacts to off-host storage

Weekly:

- restore latest backup to an isolated server
- run Odoo with `--stop-after-init`
- verify login, attachments, and key reports

DB16 · PITFALLS

Common mistakes

- Killing sessions without checking whether they are blockers or victims.
- Leaving psql transactions open with state idle in transaction.
- Refreshing staging from production while Odoo cron jobs are active.
- Backing up databases and filestores on different schedules so they no longer match.
- Keeping no written RPO, RTO, retention, or restore procedure.

Remember: Validate fixes in a disposable database before production.

DB16 · EXTRA

Senior rule

Senior rule: During an incident, identify the blocking session and business impact before terminating anything.

SECTION WRAP

You should now be able to...

- DB01 Install PostgreSQL on Ubuntu
- DB02 psql CLI, Users, Roles, and Grants
- DB03 Create and Manage Databases
- DB04 PostgreSQL Data Types for Odoo Developers
- DB05 DDL: CREATE, ALTER, and DROP
- DB06 Primary Keys, Foreign Keys, and Constraints
- DB07 CRUD: INSERT, SELECT, UPDATE, and DELETE
- DB08 Filtering with WHERE, LIKE, ILIKE, IN, and BETWEEN
- ... and 8 more lessons in this deck

Remember: Complete outstanding labs before starting the next section.

NEXT STEP

Complete the section labs
Then open the next presentation deck